

Lab NO: 2

Title: Linear Regression

Objective: To understand and apply the complete Machine Learning (ML) pipeline using Linear Regression for a regression (price prediction) problem.

Submitted By : Lijima Shrestha (ACE079BCT035)

Background: AI, ML, DL, and Data Science

- **Artificial Intelligence (AI):** A broad field focused on making machines perform intelligent tasks like reasoning, learning, and decision-making.
- **Machine Learning (ML):** A subset of AI where models learn patterns from data to make predictions.
- **Deep Learning (DL):** A subset of ML using multi-layer neural networks for complex patterns (images, text, speech).
- **Data Science:** Working with data end-to-end: collection, cleaning, analysis, modeling, and decision making.

Used Dependencies

- **NumPy:** Numerical operations.
- **Pandas:** Loading and preprocessing the dataset.
- **Matplotlib:** Visualization (scatter plots, regression line).
- **Scikit-learn:** Train/test split, preprocessing, Linear Regression, evaluation metrics.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
```

Task 1: Simple Linear Regression (Single Feature)

Goal: Build a linear regression model using only `housing_median_age` to predict `median_house_value`.

Following **7-step ML pipeline** as required in the assignment.

1 Data Retrieval and Collection

We load the dataset (`housing.csv`) and check:

- Shape (rows, columns)
- Column names
- Basic information

```
df = pd.read_csv("housing.csv")
df.shape, df.columns

((20640, 10),
 Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
        'total_bedrooms', 'population', 'households', 'median_income',
        'median_house_value', 'ocean_proximity'],
       dtype='object'))

df.info()

<class 'pandas.core.frame.DataFrame'>
, RangeIndex: 20640 entries, 0 to 20639
, Data columns (total 10 columns):
, #   Column            Non-Null Count  Dtype  
, ---  -- 
, 0   longitude         20640 non-null   float64
, 1   latitude          20640 non-null   float64
, 2   housing_median_age 20640 non-null   float64
, 3   total_rooms        20640 non-null   float64
, 4   total_bedrooms     20433 non-null   float64
, 5   population         20640 non-null   float64
, 6   households         20640 non-null   float64
, 7   median_income      20640 non-null   float64
, 8   median_house_value 20640 non-null   float64
, 9   ocean_proximity    20640 non-null   object  
,dtypes: float64(9), object(1)
, memory usage: 1.6+ MB
```

2 Data Cleaning

We check for missing values.

In this dataset, **total_bedrooms** has missing values, so we fill it using the **median** of that column.

- **Why median?** It is robust to outliers (unlike mean).

```
df.isna().sum().sort_values(ascending=False)

total_bedrooms      207
longitude           0
latitude            0
housing_median_age  0
total_rooms          0
population          0
households          0
median_income        0
median_house_value   0
ocean_proximity     0
dtype: int64

median_bedrooms = df["total_bedrooms"].median()
df["total_bedrooms"] = df["total_bedrooms"].fillna(median_bedrooms)

df.isna().sum().sort_values(ascending=False).head()

longitude           0
latitude            0
housing_median_age  0
total_rooms          0
total_bedrooms       0
dtype: int64
```

3 Feature Design

For **simple linear regression**, we use:

- Feature (**X**): `housing_median_age`
- Label (**y**): `median_house_value`

Why choose `housing_median_age`?

- It is a numeric feature that may have some relationship with price.
- Single-feature models are easy to interpret (slope + intercept).

```
X = df[["housing_median_age"]]    # feature
y = df["median_house_value"]       # label

X.head(), y.head()
```

```
(    housing_median_age
0              41.0
1              21.0
2              52.0
3              52.0
4              52.0,
0      452600.0
1      358500.0
2      352100.0
3      341300.0
4      342200.0
Name: median_house_value, dtype: float64)
```

4 Algorithm Selection

We select **Linear Regression** because:

- The output (house price) is continuous.
- We want a simple baseline relationship between one feature and the label.

5 Loss Function Selection

We use **Mean Squared Error (MSE)** as the loss function.

- MSE measures the average squared difference between actual and predicted values.
- Lower MSE means better prediction.

6 Model Learning (Training)

We split the dataset:

- **80% training**
- **20% testing**

Then train the Linear Regression model on the training set.

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

model1 = LinearRegression()
model1.fit(X_train, y_train)

model1
LinearRegression()
```

7 Model Evaluation

We evaluate on the test set and report:

- **MSE**
- **R² Score** (optional but useful)

R² close to 1 means better fit; close to 0 means weak explanatory power.

```
y_pred = model1.predict(X_test)

mse1 = mean_squared_error(y_test, y_pred)
r2_1 = r2_score(y_test, y_pred)

mse1, r2_1
(12939617265.100323, 0.012551235533311389)
```

□ Model Interpretation

For a simple linear regression:

$$\hat{y} = mx + c$$

- **Coefficient (slope, m):** change in predicted price for 1 unit increase in housing_median_age.
- **Intercept (c):** predicted price when housing_median_age = 0.

```
coef1 = model1.coef_[0]
intercept1 = model1.intercept_

coef1, intercept1
(np.float64(951.4618671495982), np.float64(179975.00158647486))
```

Interpretation (Task 1):

- **Slope (coefficient) ≈ 951.46**
→ For each **+1 year** increase in housing_median_age, the model predicts price increases by about **951.46 USD** (on average, holding nothing else).
- **Intercept ≈ 179975.00**
→ Predicted price when median age is 0 years (not very realistic, but part of the line equation).

Note: The R² score is very low here, meaning *housing age alone* does not explain much of the price variation.

Extra: Visualization (Regression Line + Predicted vs Actual)

We plot: 1) Scatter plot + regression line
2) Predicted vs actual values

Assumptions of Linear Regression

- Linear relationship between X and y
- Errors are normally distributed (approximately)
- Constant variance of errors (homoscedasticity)
- Observations are independent

Task 2: Multiple Linear Regression (All Features)

Goal: Build a linear regression model using **all features** (except `median_house_value`) to predict `median_house_value`.

Following the same **7-step ML pipeline**.

1 Data Retrieval and Collection

We already loaded and cleaned the dataset in Task 1.

Now we separate:

- **X (all features)** except label
- **y (label) = median_house_value**

```
X2 = df.drop(columns=["median_house_value"])
y2 = df["median_house_value"]
```

```
X2.head(), y2.head()
```

```
(   longitude  latitude  housing_median_age  total_rooms
total_bedrooms \
  0      -122.23      37.88            41.0      880.0
129.0
  1      -122.22      37.86            21.0     7099.0
1106.0
  2      -122.24      37.85            52.0     1467.0
190.0
  3      -122.25      37.85            52.0     1274.0
235.0
  4      -122.25      37.85            52.0     1627.0
280.0
```

```
    population  households  median_income ocean_proximity
  0        322.0       126.0      8.3252      NEAR BAY
  1      2401.0       1138.0     8.3014      NEAR BAY
  2        496.0       177.0      7.2574      NEAR BAY
  3        558.0       219.0      5.6431      NEAR BAY
  4        565.0       259.0      3.8462      NEAR BAY ,
```

```
0    452600.0
1    358500.0
2    352100.0
3    341300.0
4    342200.0
Name: median_house_value, dtype: float64)
```

2 Data Cleaning

We already handled missing `total_bedrooms` using median imputation.

We still need to handle **categorical data**:

- `ocean_proximity` is categorical, so we must encode it for regression.

```
# Check missing values (again, because Task 2 uses ALL features)

df.isna().sum().sort_values(ascending=False)

longitude          0
latitude           0
housing_median_age 0
total_rooms         0
total_bedrooms      0
population          0
households          0
median_income        0
median_house_value   0
ocean_proximity      0
dtype: int64

# Handle missing values

df["total_bedrooms"] =
df["total_bedrooms"].fillna(df["total_bedrooms"].median())

# Verify after cleaning
df.isna().sum().sort_values(ascending=False).head()

longitude          0
latitude           0
housing_median_age 0
total_rooms         0
total_bedrooms      0
dtype: int64
```

3 Feature Design

- Numeric features are **scaled** using `StandardScaler` (optional but helpful).

- `ocean_proximity` is **one-hot encoded** using `OneHotEncoder`.

We drop the first category in one-hot encoding to avoid the **dummy variable trap**.

```
# Use all features except the label

X2 = df.drop(columns=["median_house_value"])
y2 = df["median_house_value"]

X2.head(), y2.head()

(   longitude  latitude  housing_median_age  total_rooms
total_bedrooms \
 0      -122.23     37.88           41.0      880.0
129.0
 1      -122.22     37.86           21.0     7099.0
1106.0
 2      -122.24     37.85           52.0     1467.0
190.0
 3      -122.25     37.85           52.0     1274.0
235.0
 4      -122.25     37.85           52.0     1627.0
280.0

    population  households  median_income  ocean_proximity
0        322.0       126.0       8.3252    NEAR BAY
1      2401.0      1138.0       8.3014    NEAR BAY
2        496.0       177.0       7.2574    NEAR BAY
3        558.0       219.0       5.6431    NEAR BAY
4        565.0       259.0       3.8462    NEAR BAY ,
0      452600.0
1      358500.0
2      352100.0
3      341300.0
4      342200.0
Name: median_house_value, dtype: float64)

# Identify numeric and categorical columns

num_cols = X2.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = [c for c in X2.columns if c not in num_cols] # e.g.,
ocean_proximity

print("Numeric columns:", num_cols)
print("Categorical columns:", cat_cols)

Numeric columns: ['longitude', 'latitude', 'housing_median_age',
'total_rooms', 'total_bedrooms', 'population', 'households',
'median_income']
,Categorical columns: ['ocean_proximity']
```

```

# Feature scaling + One-hot encoding setup

from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder

preprocess = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), num_cols), # scaling numeric
        ("cat", OneHotEncoder(handle_unknown="ignore", drop="first"),
cat_cols) # encoding categorical
    ]
)

preprocess
ColumnTransformer(transformers=[('num', StandardScaler(),
['longitude', 'latitude',
'housing_median_age',
'total_rooms', 'total_bedrooms',
'population',
'households', 'median_income']),
('cat',
OneHotEncoder(drop='first',
handle_unknown='ignore'),
['ocean_proximity'])])

```

4 Algorithm Selection

We again use **Multiple Linear Regression**:

- Output is continuous (price)
- We want to use many factors together to better predict price.

```

# Multiple Linear Regression

from sklearn.linear_model import LinearRegression

lr_model = LinearRegression()
lr_model
LinearRegression()

```

5 Loss Function Selection

We use the same loss: **Mean Squared Error (MSE)**.

```
# We use Mean Squared Error (MSE) for evaluation
```

```
from sklearn.metrics import mean_squared_error, r2_score  
mean_squared_error  
  
<function sklearn.metrics._regression.mean_squared_error(y_true,  
y_pred, *, sample_weight=None, multioutput='uniform_average')>
```

6 Model Learning (Training)

We build a pipeline: 1) Preprocess numeric + categorical features
2) Train Linear Regression model

```
# Train-test split
X2_train, X2_test, y2_train, y2_test = train_test_split(
    X2, y2, test_size=0.2, random_state=42
)

num_cols = X2.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = [c for c in X2.columns if c not in num_cols]

preprocess = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), num_cols),
        ("cat", OneHotEncoder(handle_unknown="ignore", drop="first"),
cat_cols)
    ]
)

model2 = Pipeline(steps=[
    ("preprocess", preprocess),
    ("model", LinearRegression())
])
model2.fit(X2_train, y2_train)
model2

Pipeline(steps=[('preprocess',
                 ColumnTransformer(transformers=[('num',
                                                 StandardScaler(),
                                                 ['longitude',
                                                 'latitude'],
                                                 'housing_median_age',
                                                 'total_rooms',
                                                 'total_bedrooms',
                                                 'population',
                                                 'households',
                                                 'median_income']),
                               ('cat',
                                 OneHotEncoder(drop='first'))),
                ])
```

```
handle_unknown='ignore') ,  
['ocean_proximity']))],  
('model', LinearRegression()))])
```

7 Model Evaluation

We compute:

- **MSE**
- **R² Score**

```
y2_pred = model2.predict(X2_test)  
  
mse2 = mean_squared_error(y2_test, y2_pred)  
r2_2 = r2_score(y2_test, y2_pred)  
  
mse2, r2_2  
(4908476721.156619, 0.6254240620553604)
```

□ Model Interpretation

- **Intercept:** base predicted value when all (processed) features are 0
- **Coefficients:** effect of each feature on prediction

⚠ Since numeric features are **standardized**, their coefficients represent:

change in prediction for a **1 standard deviation** increase in that feature.

For `ocean_proximity` (one-hot), coefficients represent effect compared to the **dropped baseline category**.

```
# Extract intercept and coefficients with names  
pre = model2.named_steps["preprocess"]  
feature_names = pre.get_feature_names_out()  
  
coefs = model2.named_steps["model"].coef_  
intercept = model2.named_steps["model"].intercept_  
  
intercept  
np.float64(219899.77658329843)  
  
coef_table = pd.DataFrame({  
    "Feature": feature_names,  
    "Coefficient": coefs  
})  
coef_table["abs"] = coef_table["Coefficient"].abs()
```

```

coef_table = coef_table.sort_values("abs",
ascending=False).drop(columns="abs")

coef_table

```

	Feature	Coefficient
9	cat_ocean_proximity_ISLAND	136125.072615
7	num_median_income	75167.774766
1	num_latitude	-54415.696144
0	num_longitude	-53826.648016
5	num_population	-43403.432427
4	num_total_bedrooms	43068.181842
8	cat_ocean_proximity_INLAND	-39786.656161
6	num_households	18382.196324
2	num_housing_median_age	13889.866189
3	num_total_rooms	-13094.251162
10	cat_ocean_proximity_NEAR BAY	-5136.642217
11	cat_ocean_proximity_NEAR OCEAN	3431.140073

Model Comparison (Task 1 vs Task 2)

- **Performance:** Task 2 performs much better because it uses many important predictors (especially `median_income`, location, etc.).
- **Interpretability:** Task 1 is easiest to interpret (only one slope + intercept), but it has weak predictive power.
- **Why multiple features help:** House prices depend on many factors; using only house age ignores most of them.